

Data Visualization

Graphing our Data: Graph Tables, Add Labels, Make Notes

Ciara Zogheib

Data Sciences Institute, University of Toronto

Prerequisites

- You have the code from class open in RStudio

Coming up, we will...

- Learn new geoms and alter ggplot defaults as intermediate users
- Work through practice code from Chapter 5 (*Graph Tables, Add Labels, Make Notes*) of Healy, K. (2018). *Data Visualization: A Practical Introduction*. Princeton University Press. This code will let us
 - Organize continuous variables by group or category (in boxplots or scatter plots)
 - Plot text
 - Label outliers
 - Write and draw on our plots

**Continuous variables by groups
or categories**

Continuous variables by group or category

- We will use a new sample dataset (**organdata**, from the **socviz** library) containing numerical and categorical information on organ donations in seventeen OECD countries
- We will use a pipeline to select five rows at random and view the first six columns for those rows:

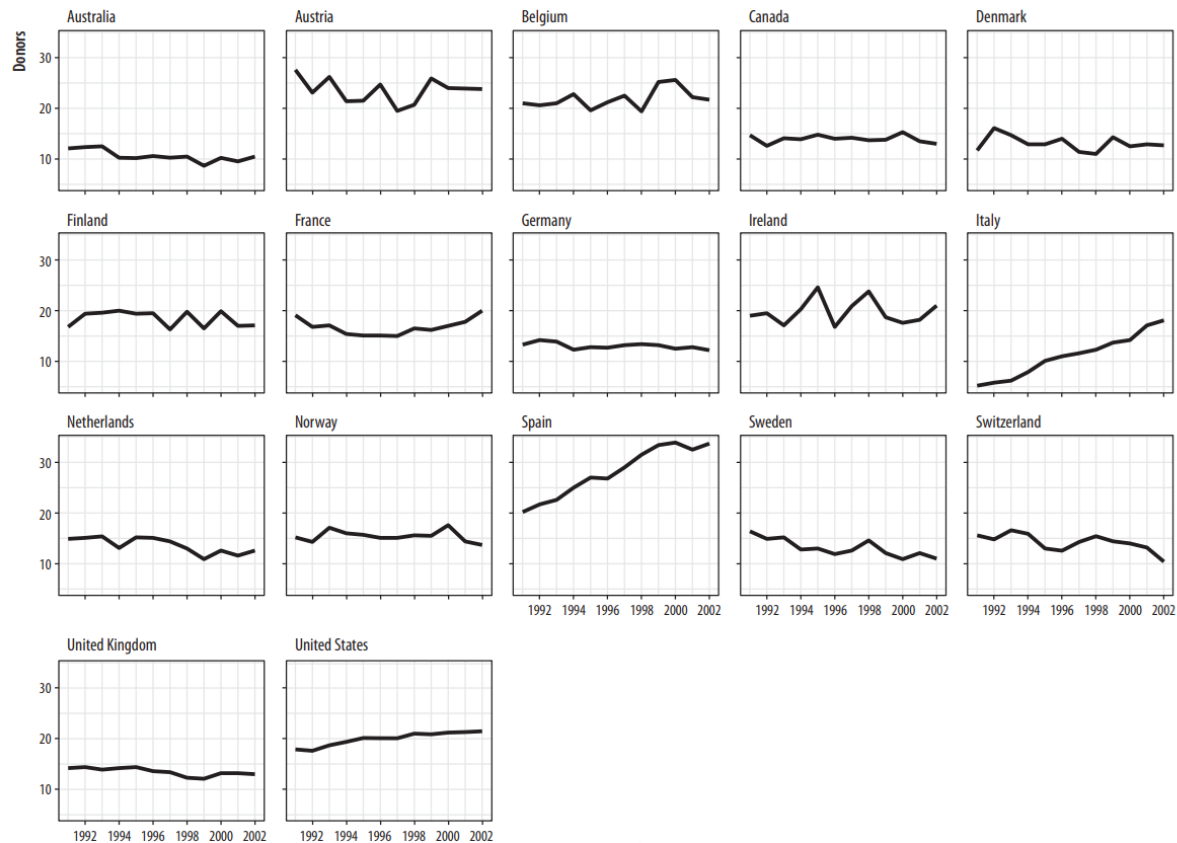
```
organdata |> select(1:6) |> sample_n(size = 10)
```

Recall: Groups and facets

- We can use `geom_line()` to explore each country's time series, incorporating grouping and faceting:

```
p <- ggplot(data = organdata, mapping = aes(x = year, y = donors))  
p + geom_line(aes(group = country)) + facet_wrap(~country)
```

Recall: Groups and facets - Result




geom_boxplot()

- We can use `geom_boxplot()` to focus on variation across countries rather than a time trend
- The `stat_boxplot()` function works in `geom_boxplot()` to calculate statistics allowing a box and whiskers to be drawn
- We tell `geom_boxplot()` the variable we want to categorize by (`country`) and the continuous variable we want to be summarized (`donors`)

```
p <- ggplot(data = organdata, mapping = aes(x = country, y =  
donors))  
  
p + geom_boxplot()
```




geom_boxplot()



```
p <- ggplot(data = organdata, mapping = aes(x = country, y = donors))  
p + geom_boxplot()
```

geom_boxplot()

By default, country names on the x axis overlap.



```
p <- ggplot(data = organdata, mapping = aes(x = country, y = donors))  
p + geom_boxplot()
```

Improving our boxplot

- We use `coord_flip()` to switch the axes and not the mappings (so our country text will not overlap)

```
p <- ggplot(data = organdata, mapping = aes(x = country, y =  
donors))  
  
p + geom_boxplot() + coord_flip()
```


Improving our boxplot

- To make the plot easier to read, we can list the countries from high to low mean donation using the `reorder()` function

```
p <- ggplot(data = organdata, mapping = aes(x = reorder(country,  
donors, na.rm = TRUE), y = donors))  
  
p + geom_boxplot() + coord_flip()
```

- The three arguments in the `reorder()` function are
 - The categorical variable we want to reorder (`country`)
 - The variable we want to reorder by (`donors`)
 - Removing missing values when calculating the mean (`na.rm=TRUE`)

Improving our boxplot - Result



```
p <- ggplot(data = organdata, mapping = aes(x = reorder(country, donors, na.rm = TRUE), y = donors))  
p + geom_boxplot() + coord_flip()
```

Activity - More grouped continuous data


- If we have a small number of observations, it can be helpful to show individual observations (points) instead of boxplots
- We can also colour our points based on another variable (`world`)

Recalling previous classes, can you modify our code (below) to plot individual points and to map the 'world' variable to colour?

```
p <- ggplot(data = organdata, mapping = aes(x = x = reorder(country,
donors, na.rm = TRUE, y = donors))

p + geom_boxplot() + coord_flip()
```

Activity - More grouped continuous data



```
p <- ggplot(data = organdata, mapping =  
  aes(x = reorder(country, donors,  
    na.rm=TRUE), y = donors, color = world))  
  
p + geom_point() + coord_flip()
```


We can see that some of our observations overlap. This makes it challenging to tell how many observations are at a given value.

Addressing overlapping observations

- By slightly perturbing our data, we can better see how many observations there are at a given value
- We do this by using `geom_jitter()` to randomly and slightly nudge each observation

```
p <- ggplot(data = organdata, mapping = aes(x = x = reorder(country,
donors, na.rm = TRUE, y = donors, color=world))
p + geom_jitter() + coord_flip()
```


Addressing overlapping observations - Result



```
p <- ggplot(data = organdata, mapping = aes(x = x = reorder(country, donors, na.rm = TRUE,  
y = donors, color=world))  
  
p + geom_jitter() + coord_flip()
```

Categorical variables with a single point

- We can take a similar approach to summarizing categorical variables with a single point per category
- The output will be a **Cleveland dotplot**
- We will use a dplyr pipeline to aggregate our organ donor data into a smaller table of summary statistics for each country

Summarizing our data

```
by_country <- organdata |> group_by(consent_law, country) |>  
  summarize_if(is.numeric, funs(mean, sd), na.rm = TRUE) |>  
  ungroup()
```


- The above code
 - Creates a new dataset called `by_country`
 - Groups our data by `consent_law` and `country`
 - Uses the `summarize_if()` function to create new variables that are the mean or standard deviation of numeric variables from our original dataset
 - Ungroups the data so our result is a tibble

Categorical variables with a single point

- Now that we have summarized our data, we can make our Cleveland dotplot using `geom_point()`
- We will also
 - Colour our results by the consent law for each country
 - Move our legend to the top of our plot
 - Add axis labels

```
p <- ggplot(data = by_country, mapping = aes(x = donors_mean,  
y = reorder(country, donors_mean), color = consent_law))  
  
p + geom_point(size=3) + labs(x = "Donor Procurement Rate", y  
= "", color = "Consent Law") + theme(legend.position="top")
```

Categorical variables with a single point - Result



```
p <- ggplot(data = by_country, mapping = aes(x = donors_mean, y = reorder(country, donors_mean),  
color = consent_law))
```

```
p + geom_point(size=3) + labs(x = "Donor Procurement Rate", y = "", color = "Consent Law") +  
theme(legend.position="top")
```


Plotting text

Data labels

- Adding labels to the points in our scatterplots can help to make our plots more informative
- We can add text to our plots using `geom_text()`

```
p <- ggplot(data = by_country, mapping = aes(x = roads_mean,  
y = donors_mean))  
p + geom_point() + geom_text(mapping = aes(label = country))
```


Data labels - Result



```
p <- ggplot(data = by_country, mapping = aes(x = roads_mean, y = donors_mean))  
p + geom_point() + geom_text(mapping = aes(label = country))
```


Data labels - Result

The text in our image is plotted directly on top of the points because both use the same x and y mapping.




```
p <- ggplot(data = by_country, mapping = aes(x = roads_mean, y = donors_mean))  
p + geom_point() + geom_text(mapping = aes(label = country))
```

Positioning our text

- We can address this problem by removing `geom_point()` (so we only see the text) OR by adjusting the position of our text
- We move our data labels by adding the `hjust` argument to `geom_text()`
 - `hjust = 0` left-justifies our labels
 - `hjust = 1` right-justifies our labels

```
p <- ggplot(data = by_country, mapping = aes(x = roads_mean,  
y = donors_mean))  
  
p + geom_point() + geom_text(mapping = aes(label = country),  
hjust=0)
```

Positioning our text - Result



```
p <- ggplot(data = by_country, mapping = aes(x = roads_mean, y =  
donors_mean))  
  
p + geom_point() + geom_text(mapping = aes(label = country), hjust=0)
```

ggrepel

- For more flexible text formatting, install and load the `ggrepel` library

```
install.packages("ggrepel")  
  
library(ggrepel)
```

- To explore what we can do with `ggrepel`, we will use a new dataset about historical United States presidents (`elections_historic`) from the `socviz` library we previously loaded

ggrepel

```
p <- ggplot(elections_historic, aes(x = popular_pct, y = ec_pct,  
  label = winner_label))  
  
p + geom_hline(yintercept = 0.5, size = 1.4, color = "gray80") +  
  geom_vline(xintercept = 0.5, size = 1.4, color = "gray80") +  
  geom_point() +  
  geom_text_repel() +  
  scale_x_continuous(labels = scales::percent) +  
  scale_y_continuous(labels = scales::percent) +  
  
labs(x = "Winner's share of Popular Vote", y = "Winner's share of Electoral  
College Votes", title = "Presidential Elections: Popular & Electoral  
College Margins", subtitle = "1824-2016")
```

ggrepel

```
p <- ggplot(elections_historic, aes(x = popular_pct, y = ec_pct,  
  label = winner_label))  
  
p + geom_hline(yintercept = 0.5, size = 1.4, color = "gray80") +  
  geom_vline(xintercept = 0.5, size = 1.4, color = "gray80") +  
  geom_point() +  
  geom_text_repel() +  
  scale_x_continuous(labels = scales::percent)  
  scale_y_continuous(labels = scales::percent)  
  
labs(x = "Winner's share of Popular Vote", y =  
  "College Votes", title = "Presidential Elections: Popular & Electoral  
  College Margins", subtitle = "1824-2016")
```

Vote shares are stored as proportions rather than percents, so we adjust the labels of the scales.

ggrepel

```
p <- ggplot(elections_historic, aes(x = popular_pct, y = ec_pct,  
  label = winner_label))  
  
p + geom_hline(yintercept = 0.5, size = 1.4, color = "gray80") +  
  geom_vline(xintercept = 0.5, size = 1.4, color = "gray80") +  
  geom_point() +  
  geom_text_repel() +  
  scale_x_continuous(labels = scales::percent) +  
  scale_y_continuous(labels = scales::percent) +  
  labs(x = "Winner's share of Popular Vote", y = "College Votes", title = "Presidential Elections: Popular & Electoral  
  College Margins", subtitle = "1824-2016")
```


**We add reference lines so
we can see the 50%
threshold of votes on
each axis.**

ggrepel

```
p <- ggplot(elections_historic, aes(x = popular_pct, y = ec_pct,  
  label = winner_label))  
  
p + geom_hline(yintercept = 0.5, size = 1.4, color = "gray80") +  
  geom_vline(xintercept = 0.5, size = 1.4, color = "gray80") +  
  geom_point() +  
  geom_text_repel() +  
  scale_x_continuous(labels = scales::percent) +  
  scale_y_continuous(labels = scales::percent) +  
  
labs(x = "Winner's share of Popular Vote", y = "Winner's share of Electoral  
College Votes", title = "Presidential Elections: Popular & Electoral  
College Margins", subtitle = "1824-2016")
```

**geom_text_repel() will
ensure that our data
labels do not overlap.**

ggrepel - Result



Labeling outliers

Labeling specific points


- Sometimes, we only want to label specific points of interest (such as outliers) rather than every single item
- We do this by using the `subset()` function to tell `geom_text_repel()` to use a different dataset from the one being used by `geom_point()`
- We will return to our `by_country` dataset to explore labeling outliers

Labeling specific points

- In the below code, we use `subset()` to select only the cases from our dataset (`by_country`) where `gdp_mean` is over 25000
- Only points within this subset are labeled

```
p <- ggplot(data = by_country, mapping = aes(x = gdp_mean, y =  
health_mean))  
  
p + geom_point() +  
  geom_text_repel(data = subset(by_country, gdp_mean > 25000),  
mapping = aes(label = country))
```

Labeling specific points - Result



Writing and drawing in our plots

Annotating plots

- It can sometimes be useful to annotate a figure or place arbitrary text on a plot
 - That is, text that is not mapped to a variable
- We do this using the `annotate()` function, which is not a geom, but uses the features of geoms (such as the ability to modify size, colour, and x or y position)

Annotating plots with text

```
p <- ggplot(data = organdata, mapping = aes(x = roads, y = donors))  
p + geom_point() + annotate(geom = "text", x = 91, y = 33, label =  
"A surprisingly high \n recovery rate.", hjust = 0)
```


- The above code
 - Uses `annotate()` to add text
 - Positions the text at (91, 33) on our plot
 - Using the **newline code** (`'\n'`) to force a line break
 - Left-justifies our text

Annotating plots with shapes

- We can also use `annotate` to add shapes to our plot
- For example, we can add a red rectangle to our previous code by calling `annotate()` again with “rect” instead of “text”

```
p <- ggplot(data = organdata, mapping = aes(x = roads, y = donors))  
p + geom_point() +  
  annotate(geom = "rect", xmin = 85, xmax = 135, ymin = 30, ymax = 35,  
    fill = "red", alpha = 0.2) +  
  annotate(geom = "text", x = 91, y = 33, label = "A surprisingly high  
  \n recovery rate.", hjust = 0)
```

Annotating plots - Result



Next...

- Visualization with Purpose
- Refining our ggplots in R, with consideration for aesthetic qualities of data visualization and colour theory
- Strategies for accessible data visualization (Ethics and inequity)
- Exploring data visualization as a tool for advocacy (Inequity)